

# **IMPERIAL**

## **Quantitative Risk Management Coursework 1**

Imperial College London

Yve Meng (06060863)

Harry Qiao (02057393)

Yibo Wang (02199425)

Ganwen Zuo (02222591)

MSc Mathematics & Finance

November 20th 2025

# Contents

<b>1 Part A: Stylised Facts and GARCH Modelling</b>	<b>2</b>
1.1 Introduction . . . . .	2
1.2 Stylised Facts . . . . .	2
1.3 GARCH(1,1) Model with Constant Mean and N(0,1) Innovations . . . . .	3
1.4 ARMA(1,1)–GARCH(1,1) Model with Student-t Innovations . . . . .	6
<b>2 Part B: Risk Measures</b>	<b>10</b>
2.1 Introduction . . . . .	10
2.2 Historical Simulation (HS) . . . . .	10
2.3 FHS with EWMA . . . . .	10
2.4 FHS with GARCH . . . . .	12
2.5 Model Backtest and Recommendation . . . . .	13
<b>A Additional Tables and Figures</b>	<b>16</b>
<b>B Code</b>	<b>17</b>

# 1 Part A: Stylised Facts and GARCH Modelling

## 1.1 Introduction

This report analyses the daily log returns of the EURO STOXX 50 index (Bloomberg, 2025) with the aim of identifying key stylised facts, modelling log returns using GARCH processes, and evaluating the adequacy of different conditional mean and innovation specifications.

The dataset consists of daily closing values of the EURO STOXX 50 (ticker: SX5E) from 30 October 2013 to 27 October 2023. The EURO STOXX 50 is a benchmark index tracking the fifty largest and most liquid blue-chip companies in the Eurozone, providing a representative measure of regional economic conditions (Bloomberg, 2025).

With one observation per trading day, the dataset contains 2,512 data points. Our analysis uses logarithmic daily returns, which are more appropriate for time-series modelling and volatility analysis. Let  $S_t$  denote the closing value on day  $t$ ; the log-return (in percentage) is computed as:

$$r_t := (\log(S_t) - \log(S_{t-1})) \times 100. \quad (1)$$

## 1.2 Stylised Facts

To characterise the distributional properties of the log returns, we compute their sample mean, standard deviation, skewness, and kurtosis.

**Table 1:**  
*Summary statistics of log returns (%)*

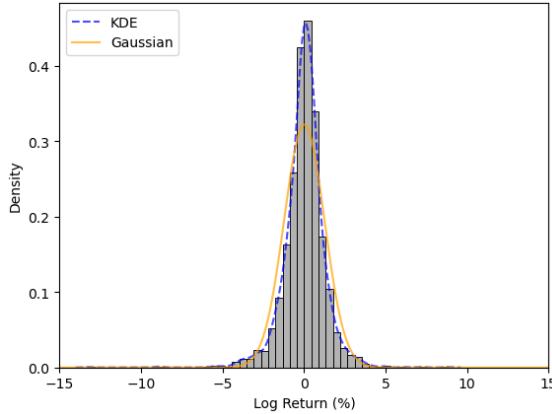
Mean	Standard Deviation	Skewness	Kurtosis
0.0118	1.24	-0.796	10.3

As shown in Table 1, the log returns (in percentage) have a small but positive mean, indicating modest average growth over the sample period. The standard deviation is large relative to the mean, highlighting substantial day-to-day variability. The return distribution displays clear departures from normality. The skewness of  $-0.796$  indicates a longer, heavier left tail, while the kurtosis of 10.3 (excess kurtosis 7.3) points to a strongly leptokurtic distribution with a high peak and heavy tails.

Figure 1 further illustrates the deviation of log returns from normality. The kernel density estimate (KDE) in blue shows a sharper peak and heavier tails than the fitted normal density (orange) based on the sample mean and variance. This discrepancy is consistent with the summary statistics in Table 1.

To assess serial dependence in the log returns, we compute the empirical autocorrelation functions (ACFs) of the log returns, their absolute values, and their squares. These diagnostics capture linear dependence as well as volatility clustering and persistence. As shown in Figure 2(a), the ACF of the log returns exhibits no significant autocorrelation,

**Figure 1:**  
*Histogram and KDE of the log returns*



indicating that log returns are serially uncorrelated. In contrast, Figures 2(b) and (c) show clear positive autocorrelations in the absolute and squared log returns, reflecting the volatility clustering effect in which large movements tend to be followed by further large movements.

### 1.3 GARCH(1,1) Model with Constant Mean and N(0,1) Innovations

To model the log returns, we begin by fitting a GARCH(1,1) model with a constant conditional mean and standard normal innovations. The model is specified as follows:

$$X_t = \mu + \sigma_t Z_t, \quad (2)$$

$$\sigma_t^2 = \omega + \alpha_1(X_{t-1} - \mu)^2 + \beta_1\sigma_{t-1}^2, \quad t \in \mathbb{Z}. \quad (3)$$

The innovations  $Z_t$  are assumed to be i.i.d.  $N(0,1)$  random variables. The model is fitted using the `arch` package in Python. For the fitted model, the maximised log-likelihood is -3738.92, with AIC = 7485.84 and BIC = 7509.14.

**Table 2:**  
*Coefficient outputs of the GARCH(1,1) model with constant mean and N(0,1) innovations*

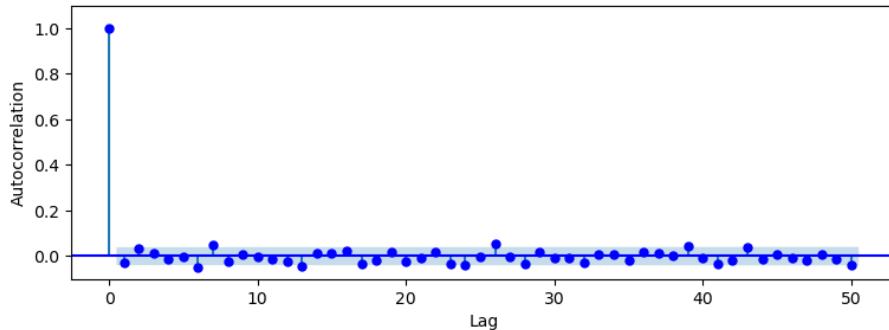
Variable	Estimate	Std. Error	T Value	$\mathbb{P} >  t $
$\mu$	0.0476	$2.02 \times 10^{-2}$	2.36	$1.83 \times 10^{-2}$
$\omega$	0.0596	$1.91 \times 10^{-2}$	3.13	$1.77 \times 10^{-3}$
$\alpha_1$	0.139	$3.06 \times 10^{-2}$	4.55	$5.35 \times 10^{-6}$
$\beta_1$	0.824	$3.46 \times 10^{-2}$	23.8	$2.36 \times 10^{-125}$

All coefficient estimates reported in Table 2 are statistically significant at the 5% level. The estimated parameters,  $\alpha_1 = 0.139$  (95% CI:  $[7.91 \times 10^{-2}, 0.199]$ ) and  $\beta_1 = 0.824$  (95% CI:  $[0.756, 0.892]$ ), sum to 0.963, which is close to but strictly below 1. This high degree of persistence is typical for equity index returns and reflects the slow decay of

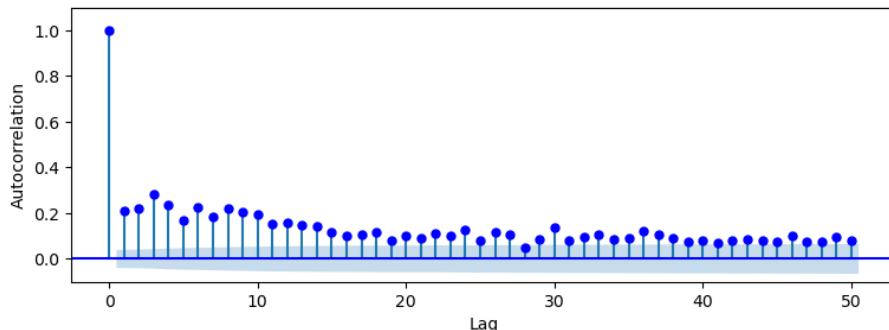
**Figure 2:**

*Empirical ACFs of the log returns, their absolute values, and their squares*

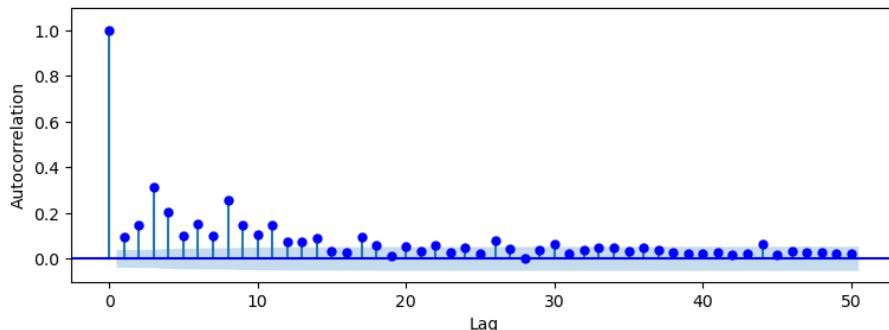
(a) Empirical ACF of the log returns



(b) Empirical ACF of the absolute log returns



(c) Empirical ACF of the squared log returns

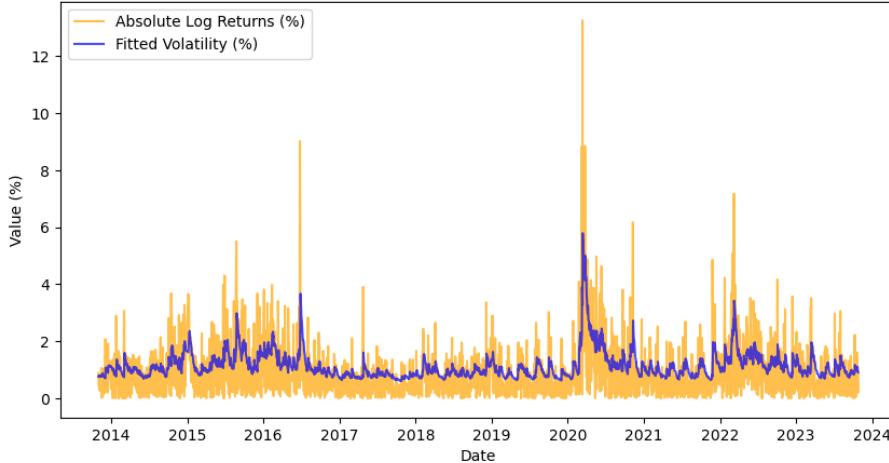


volatility shocks over time. The constant term in the variance equation  $\omega = 0.0596$  (95% CI:  $[2.23 \times 10^{-2}, 9.70 \times 10^{-2}]$ ) is small relative to the size of the volatility process. The conditional mean estimate,  $\mu = 0.0476$  (95% CI:  $[8.07 \times 10^{-3}, 8.72 \times 10^{-2}]$ ), corresponds to an average log return of 0.0476%, which is statistically significant but economically small.

Figure 3 compares the absolute log returns with the fitted volatility from the GARCH(1,1) model with constant conditional mean and standard normal innovations. The fitted volatility closely tracks periods of heightened market turbulence, rising sharply during times of

**Figure 3:**

Absolute log returns and fitted volatility of the GARCH(1,1) model with constant mean and  $N(0,1)$  innovations



large absolute returns (such as the volatility spikes in 2016 and early 2020). This shows that the model captures volatility clustering and the persistence of high-variance regimes. Although it smooths the most extreme movements, the fitted volatility still reflects the overall pattern of market variability over time.

To assess the model's goodness of fit, we examine whether the standardised residuals follow an i.i.d. sample from the assumed  $N(0,1)$  distribution. Moreover, a correctly specified model should yield residuals with no significant serial correlation. We evaluate this by analysing the ACFs of the standardised residuals, plotting their histogram and KDE, inspecting the normal QQ-plot, and performing the one-sample Anderson–Darling test (Nelson, 1998).

The empirical ACFs in Figure 4 shows that the standardised residuals display no meaningful serial correlation, which is consistent with the i.i.d. assumption required for a correctly specified model.

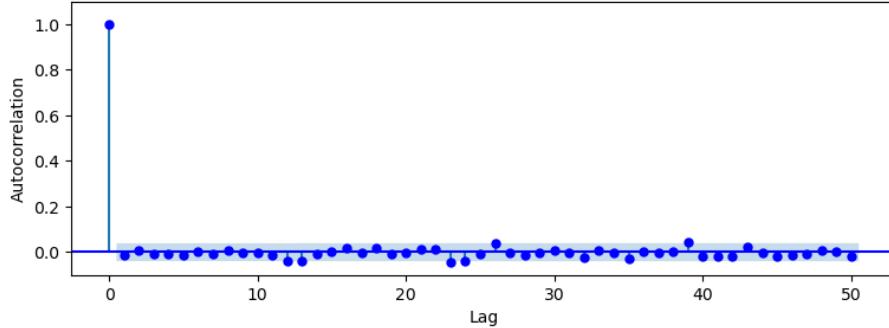
The histogram and KDE of the standardised residuals in Figure 5(a) indicate clear departures from the  $N(0,1)$  distribution as the empirical density exhibits higher peak heavier tails than that of the  $N(0,1)$ . This observation is reinforced by the QQ-plot in Figure 5(b), where the lower and upper sample quantiles deviate substantially from the reference line  $y = x$ . The lower tail falls below the theoretical quantiles, while the upper tail lies above them, indicating heavier tails than  $N(0,1)$ . These plots suggest that the  $N(0,1)$  innovation assumption is not adequate and that the model may not be correctly specified.

The one-sample Anderson–Darling test provides strong evidence against normality. The test statistic is 14.1, which exceeds the 1% critical value of 1.09. Consequently, we reject the null hypothesis that the standardised residuals follow a  $N(0,1)$  distribution.

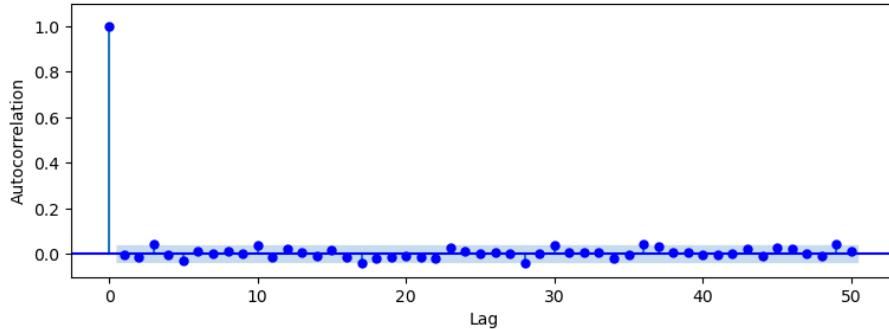
**Figure 4:**

Empirical ACFs of the standardised residuals and their absolute values of the GARCH(1,1) model with constant mean and  $N(0,1)$  innovations

(a) Empirical ACF of the standardised residuals



(b) Empirical ACF of the absolute standardised residuals



## 1.4 ARMA(1,1)–GARCH(1,1) Model with Student-t Innovations

To allow for both serial dependence in the conditional mean and heavier-tailed innovations, we replace the constant conditional mean with an ARMA(1,1) process and assume Student-t distributed innovations. The model is specified as follows:

$$X_t = \mu_t + \sigma_t Z_t, \quad (4)$$

$$\mu_t = \mu + \phi_1(X_{t-1} - \mu) + \theta_1(X_{t-1} - \mu_{t-1}), \quad (5)$$

$$\sigma_t^2 = \alpha_0 + \alpha_1(X_{t-1} - \mu_{t-1})^2 + \beta_1\sigma_{t-1}^2, \quad t \in \mathbb{Z}. \quad (6)$$

The innovations  $Z_t$  follow a normalised Student-t distribution with degree of freedom  $\nu$ . Because the `arch` package in Python does not allow ARMA(1,1)–GARCH(1,1) to be fitted directly, the parameters are estimated using a constrained numerical optimiser that maximises the log-likelihood function:

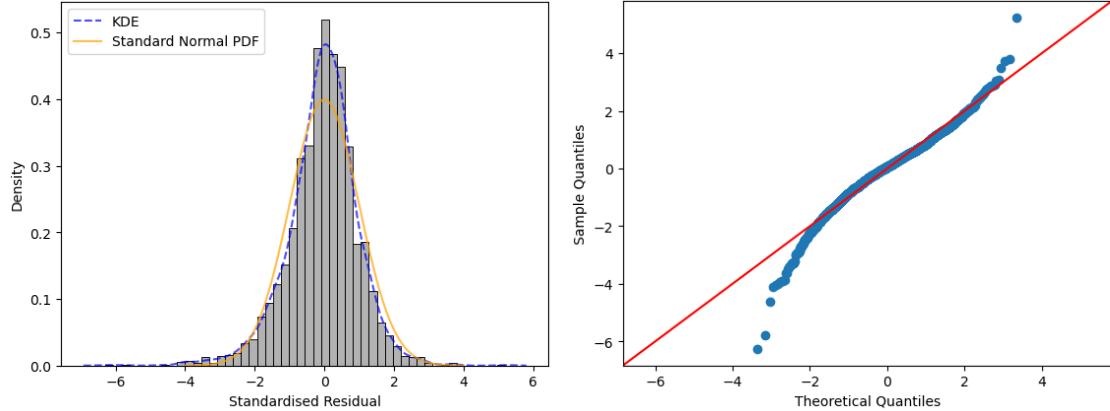
$$\ell(\alpha_0, \alpha_1, \beta_1, \mu, \phi_1, \theta_1, \nu; X) = \sum_{t=1}^T \left[ -\log \sigma_t + \log f_Z \left( \frac{X_t - \mu_t}{\sigma_t} \right) \right]. \quad (7)$$

To ensure stable convergence, the initial parameter values are set using a combination of empirical moments and estimates from an auxiliary ARMA(1,1) model. The GARCH

**Figure 5:**

Diagnostic plots of standardised residuals of the GARCH(1,1) model with constant mean and  $N(0,1)$  innovations

(a) Histogram of the standardised residuals      (b) QQ-plot of standardised residuals



constant  $\alpha_0$  is initialised using the identity  $\alpha_0 = \text{Var}(X_t)(1 - \alpha_1 - \beta_1)$  with choices  $\alpha_1 = 0.1$  and  $\beta_1 = 0.85$ . Starting values for the ARMA parameters  $\phi_1$  and  $\theta_1$  come from an ARMA(1,1) model fitted via `statsmodels`, while  $\mu$  is set to the sample mean of the log returns. The degrees-of-freedom parameter  $\nu$  is initialised at 10 to provide a suitably heavy-tailed starting point. Optimisation is performed using the SLSQP algorithm subject to parameter bounds and the stationarity constraint  $\alpha_1 + \beta_1 < 1$ .

**Table 3:**

Coefficient estimates of the ARMA(1,1)–GARCH(1,1) model with Student- $t$  innovations

Variable	Estimate	Std. Error	T Value	$\mathbb{P} >  t $
$\alpha_0$	0.0475	$1.28 \times 10^{-2}$	3.71	$2.17 \times 10^{-4}$
$\alpha_1$	0.148	$2.42 \times 10^{-2}$	6.11	$1.03 \times 10^{-9}$
$\beta_1$	0.834	$2.44 \times 10^{-2}$	34.1	$1.52 \times 10^{-255}$
$\mu$	0.0699	$1.35 \times 10^{-2}$	5.17	$2.39 \times 10^{-7}$
$\phi_1$	0.911	$6.88 \times 10^{-2}$	13.2	$5.13 \times 10^{-40}$
$\theta_1$	-0.934	$5.97 \times 10^{-2}$	-15.6	$3.95 \times 10^{-55}$
$\nu$	4.73	0.470	10.1	$8.30 \times 10^{-24}$

The optimiser converges successfully, with parameter estimates reported in Table 3. The model exhibits strong persistence in volatility, with  $\alpha_1 = 0.148$  (95% CI: [0.100, 0.195]),  $\beta_1 = 0.834$  (95% CI: [0.786, 0.881]), and heavy-tailed innovations, with an estimated degrees-of-freedom parameter  $\nu = 4.73$  (95% CI: [3.81, 5.65]). The ARMA(1,1) component captures short-term dependence in the conditional mean, with coefficient estimates  $\phi_1 = 0.91$  (95% CI: [0.777, 1.046]) and  $\theta_1 = -0.93$  (95% CI: [-1.05, -0.817]).

For the fitted ARMA(1,1)–GARCH(1,1) model with Student- $t$  innovations, the maximised log-likelihood is -3635.09, with AIC = 7284.18 and BIC = 7324.97. Compared

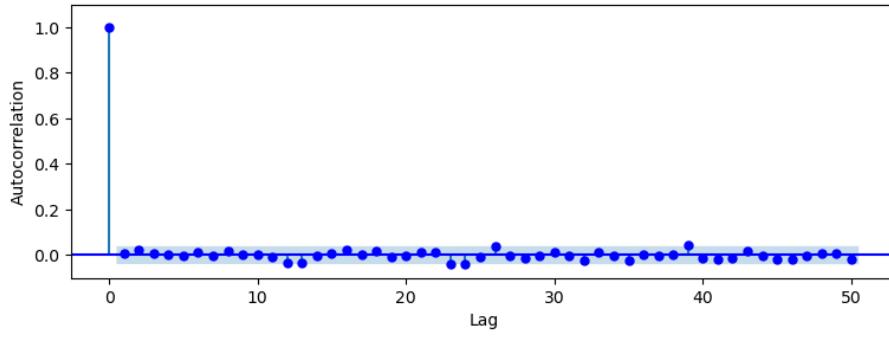
with the GARCH(1,1) model with  $N(0,1)$  innovations, the ARMA(1,1)–GARCH(1,1) model achieves a higher log likelihood and lower information criteria despite its larger number of parameters. This indicates an improvement in overall model fit.

To assess the goodness of fit, we examine whether the standardised residuals behave like an i.i.d. sample from the fitted  $t_{\hat{\nu}}$  distribution. The conditional mean and variance, and hence the standardised residuals, are computed iteratively using the sample mean and variance as initial values. A correctly specified model should also produce residuals with no meaningful serial correlation. We evaluate this by analysing the ACFs of the standardised residuals, plotting their histogram and KDE with the fitted t-density, inspecting the QQ-plot, and formally testing distributional adequacy using the k-sample Anderson–Darling test (Scholz & Stephens, 1987).

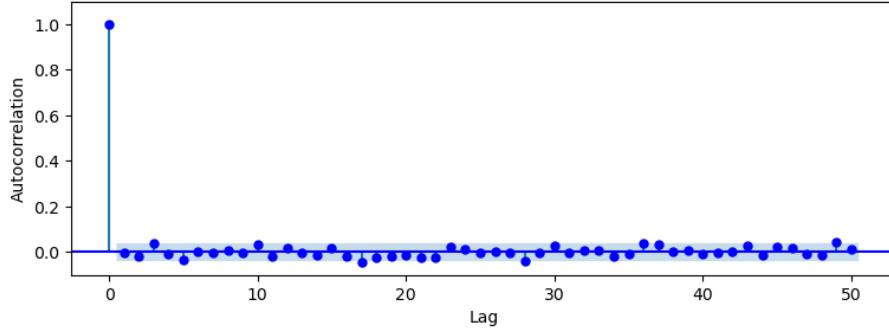
**Figure 6:**

*Empirical ACFs of the standardised residuals and their absolute values of the ARMA(1,1)–GARCH(1,1) model with Student-t innovations*

(a) Empirical ACF of the standardised residuals



(b) Empirical ACF of the absolute standardised residuals



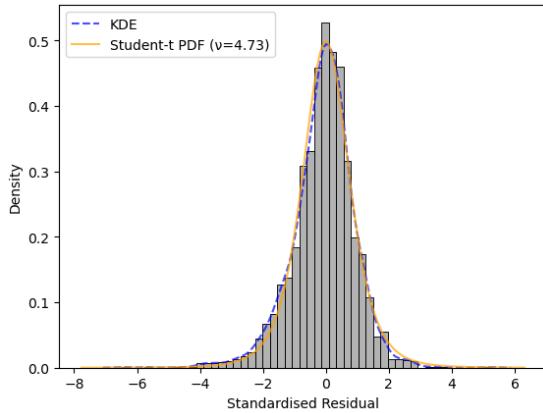
The empirical ACFs in Figure 6 shows that the standardised residuals display no meaningful serial correlation, which is consistent with the i.i.d. assumption required for a correctly specified model.

Figure 7(a) presents the histogram and KDE of the standardised residuals, overlaid with the fitted Student-t density with  $\hat{\nu}$  degrees of freedom (scaled to have variance 1). The

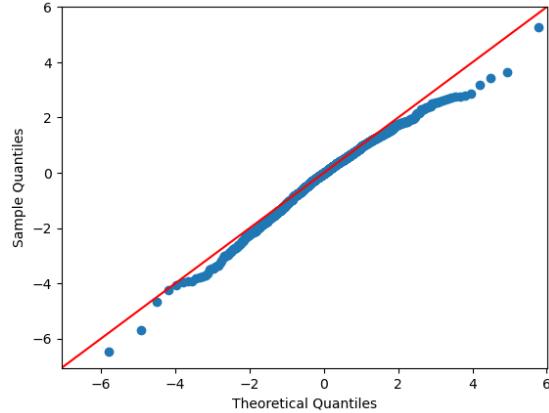
**Figure 7:**

Diagnostic plots of standardised residuals of the ARMA(1,1)–GARCH(1,1) model with Student-t innovations

(a) Histogram of the standardised residuals



(b) QQ-plot of standardised residuals



empirical distribution aligns closely with the fitted t-density, providing support for the heavy-tailed innovation assumption. Figure 7(b) shows some deviations from the Student-t distribution, though these are substantially smaller than under the normality assumption.

The k-sample Anderson–Darling test provides further evidence of remaining misspecification: the test statistic is 3.75 with an associated significance level of 0.01, leading to rejection of the null hypothesis that the standardised residuals follow a  $t_{\hat{\nu}}$  distribution. While the Student-t specification substantially improves the fit relative to the Gaussian model, these diagnostics suggest that the innovation distribution is still not fully captured by the fitted model.

To assess whether the ARMA(1,1) component is necessary, we compare the ARMA(1,1)–GARCH(1,1) model with Student-t innovations to the GARCH(1,1) model with constant conditional mean and Student-t innovations. The model without the ARMA component achieves a maximised log-likelihood of -3639.90, only marginally below the value -3635.09 obtained under the ARMA(1,1)–GARCH(1,1) model. The resulting AIC and BIC values (AIC=7289.79, BIC=7318.92) are very close to those of the more complex model, indicating that the ARMA component offers little improvement.

The coefficient outputs and diagnostic checks for the GARCH(1,1) model with constant conditional mean and Student-t innovations are reported in the Appendix (Table 6, Figure 13 and 14). There are no evidence of serial correlation in the standardised residuals, and the distributional diagnostics (histogram, KDE, QQ-plot, and k-sample Anderson–Darling test) closely mirror those of the ARMA(1,1)–GARCH(1,1) model, with the k-sample Anderson–Darling test yielding a statistic of 2.30 and significance level 0.037. These results suggest that the ARMA(1,1) component is not necessary for modelling the conditional mean of the log returns, and that a constant-mean GARCH(1,1) model with Student-t innovations provides an adequate and parsimonious specification.

## 2 Part B: Risk Measures

### 2.1 Introduction

This section evaluates one-day ahead VaR and ES forecasts for the ProShares Short Dow30 (DOG) ETF (ProShares, 2025). The daily linearised loss relative to the investment value is defined as:

$$L_{t+1}^\Delta = \frac{L_{t+1}^\Delta}{V_t} = -r_{t+1} \quad (8)$$

where  $V_t$  is the investment value and  $-r_{t+1}$  is the negative log return. We compare Historical Simulation (HS), Filtered Historical Simulation with EWMA (FHS-EWMA), and FHS-GARCH(1,1) at 95% and 99% confidence levels.

### 2.2 Historical Simulation (HS)

We utilize a sliding window of  $N = 500$  days. For each forecasting date  $t$ , the dataset is  $L = \{L_{t-499}, \dots, L_t\}$ . The estimated VaR at level  $\alpha$  is the empirical quantile:

$$\widehat{\text{VaR}}_\alpha(L_t) = q_\alpha(L) = \inf \{x \in \mathbb{R} : \mathbb{P}(L \leq x) \geq \alpha\}. \quad (9)$$

The Expected Shortfall (ES) is calculated as the expected loss conditional on exceeding the VaR:

$$\widehat{\text{ES}}_\alpha(L_t) = \frac{\mathbb{E}\left[L \mathbf{1}_{\{L \geq \widehat{\text{VaR}}_\alpha(L)\}}\right]}{1 - \alpha} = \frac{1}{\sum_{j=t-499}^t \mathbf{1}_{\{L_j \geq \widehat{\text{VaR}}_\alpha(L)\}}} \sum_{j=t-499}^t L_j \mathbf{1}_{\{L_j \geq \widehat{\text{VaR}}_\alpha(L)\}}. \quad (10)$$

The HS forecasts in Figure 8 exhibit a characteristic ‘step-like’ shape. Following the 2020 volatility spike, risk estimates remain artificially elevated for nearly two years, confirming that HS is strongly backward-looking. Extreme observations remain in the 500-day window long after market conditions stabilise, leading to risk overestimation in calm periods. Conversely, during the onset of the 2020 crash, the 95% VaR frequently fell below realised losses, indicating poor responsiveness to sudden shocks.

### 2.3 FHS with EWMA

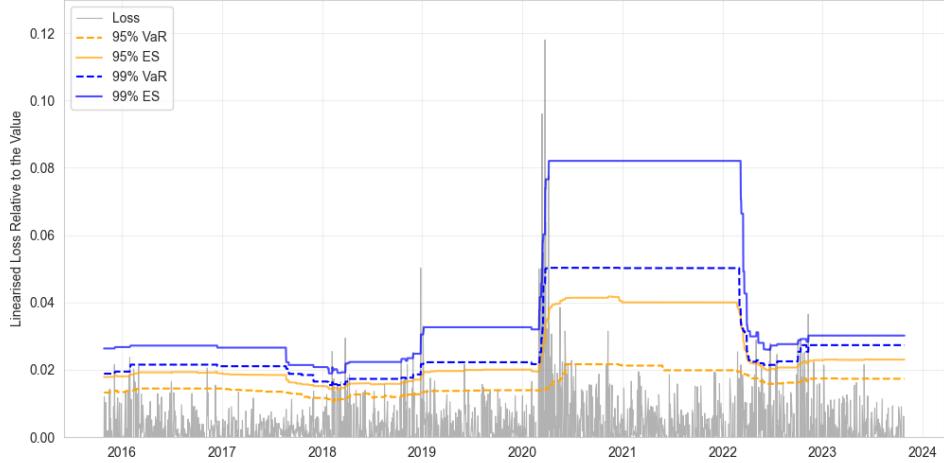
To address the lag in HS, we apply FHS using an EWMA volatility model. The volatility forecast is computed recursively:

$$\hat{\sigma}_{t+1}^2 = \alpha(L_t - \hat{\mu}_t)^2 + (1 - \alpha)\hat{\sigma}_t^2. \quad (11)$$

Using the specified parameters  $\alpha = 0.06$  and  $\hat{\mu}_t = 0$ , the update equation becomes:

$$\hat{\sigma}_{t+1}^2 = 0.06L_t^2 + 0.94\hat{\sigma}_t^2. \quad (12)$$

**Figure 8:**  
VaR and ES using Historical Simulation ( $N=500$ )



The initial volatility  $\hat{\sigma}_0$  is the variance of the first 500 observations:

$$\hat{\sigma}_0^2 = \frac{1}{n} \sum_{i=0}^{499} (L_i - \bar{L})^2. \quad (13)$$

VaR and ES are forecasted by scaling the quantile and ES of the standardised residuals ( $Z_t = L_t / \hat{\sigma}_t$ ) by the current volatility:

$$\widehat{\text{VaR}}_\alpha(L_{t+1}) = \hat{\sigma}_{t+1} q_\alpha(Z_{t+1}), \quad \widehat{\text{ES}}_\alpha(L_{t+1}) = \hat{\sigma}_{t+1} \widehat{\text{ES}}_\alpha(Z_{t+1}). \quad (14)$$

**Figure 9:**  
VaR and ES using FHS-EWMA

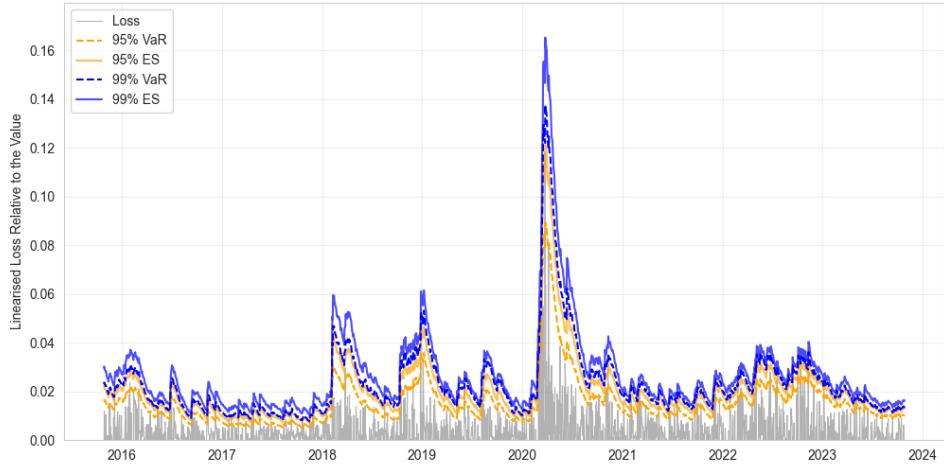
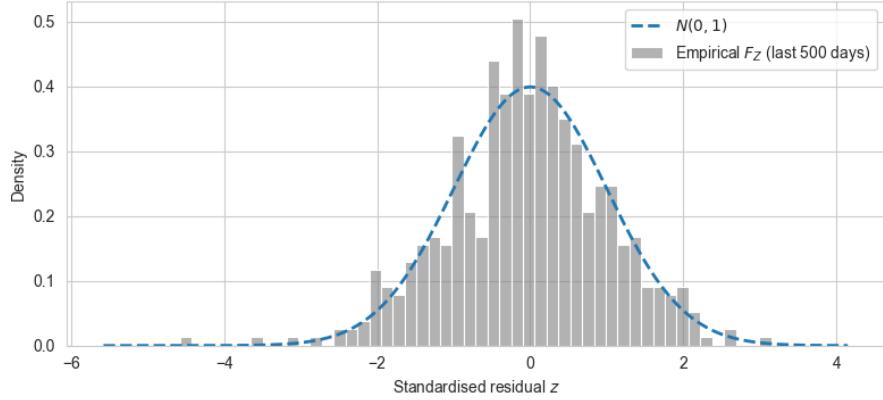


Figure 9 illustrates the results. Unlike HS, the FHS-EWMA forecasts track the realised volatility dynamically. The curves rise sharply during the 2020 shock and decay as mar-

kets stabilise, avoiding the ‘ghost effects’ of old data. The widening gap between 95% and 99% measures during stress periods reflects the heavy-tailed nature of the residuals.

**Figure 10:**  
*Empirical distribution of standardised residuals (FHS-EWMA)*



The empirical distribution of the standardised residuals (Figure 10) is roughly symmetric and bell-shaped, validating the  $\mu = 0$  assumption. However, the distribution shows higher concentration around the mean and slightly heavier tails than the standard normal, indicating excess kurtosis.

## 2.4 FHS with GARCH

We estimate VaR and ES using FHS-GARCH(1,1) with normal innovations. The process is defined as:

$$L_t = \sigma_t Z_t, \quad (15)$$

$$\sigma_t^2 = \alpha_0 + \alpha_1 L_{t-1}^2 + \beta_1 \sigma_{t-1}^2. \quad (16)$$

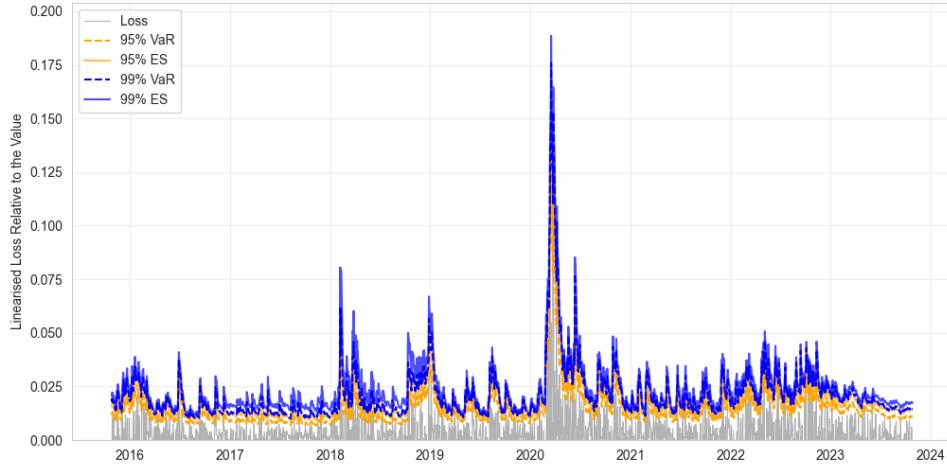
Parameters were re-estimated daily using a 500-day rolling window via conditional maximum likelihood:

$$\mathcal{L}(\alpha_0, \alpha_1, \beta_1; X, \sigma_0) = \prod_{t=1}^T \frac{1}{\sigma_t} f_Z\left(\frac{L_t}{\sigma_t}\right). \quad (17)$$

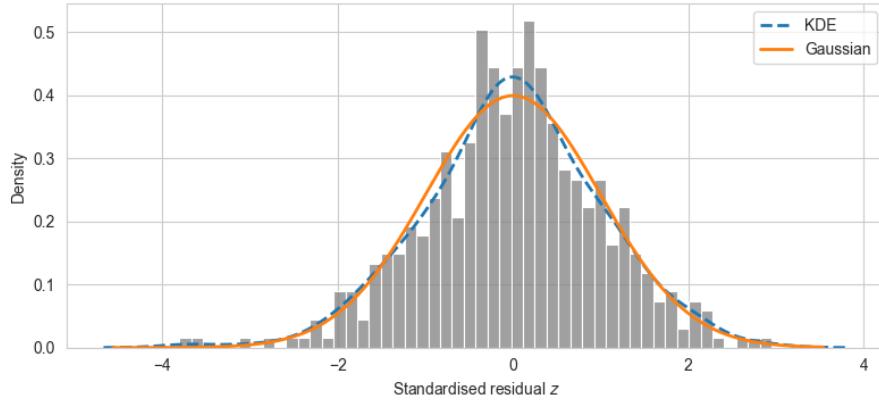
where  $f_Z$  is the standard normal density. The forecast formulas for VaR and ES remain the same as in section 3.2, utilizing the GARCH  $\hat{\sigma}_{t+1}$ .

As seen in Figure 11, the GARCH forecasts share the dynamic shape of EWMA but exhibit sharper, more frequent fluctuations. This sensitivity suggests that re-estimating GARCH parameters on a short window ( $N = 500$ ) may introduce instability. While it reacts aggressively to clustering, the GARCH model produces less conservative risk estimates during quiet periods compared to EWMA. The residual distribution (Figure 12) confirms the model successfully filters volatility clustering, showing properties similar to the EWMA residuals.

**Figure 11:**  
VaR and ES using FHS-GARCH



**Figure 12:**  
Empirical distribution of standardised residuals (FHS-GARCH)



## 2.5 Model Backtest and Recommendation

We evaluate the models using Kupiec's tests for coverage and independence (Kupiec et al., 1995). We define the violation indicator  $I_t$ :

$$I_t := \begin{cases} 1, & L_t > \text{VaR}_\alpha(L_t), \\ 0, & L_t \leq \text{VaR}_\alpha(L_t). \end{cases} \quad (18)$$

Assuming  $I_t \sim \text{Bernoulli}(\pi)$ , the likelihood under the null hypothesis ( $\pi = 1 - \alpha$ ) vs the alternative ( $\pi = \hat{\pi}$ ) is tested using the Unconditional Coverage statistic:

$$\text{LR}_{\text{uc}} := -2 \log \left( \frac{L(1 - \alpha; I)}{L(\hat{\pi}; I)} \right) \sim \chi^2(1). \quad (19)$$

For independence, we estimate transition probabilities  $\pi_{01}$  and  $\pi_{11}$  and calculate the Likelihood Ratio:

$$\text{LR}_{\text{ind}} := -2 \log \left( \frac{L(\hat{\Pi}_0; I)}{L(\hat{\Pi}_1; I)} \right) \sim \chi^2(1). \quad (20)$$

The joint test is  $\text{LR}_{cc} = \text{LR}_{uc} + \text{LR}_{\text{ind}} \sim \chi^2(2)$ . Finally, ES is backtested using the test statistic  $Z$  on the residuals  $\xi_t = (L_t - ES_\alpha)I_t$ :

$$Z := \frac{\sum_{t=1}^T \xi_t}{\sqrt{\sum_{t=1}^T \xi_t^2}} \sim N(0, 1). \quad (21)$$

**Table 4:**  
*VaR backtesting results*

<b>Model</b>	<b>Level</b>	<b>Unconditional</b>		<b>Independence</b>		<b>Joint</b>	
		$p_{uc}$	<b>Result</b>	$p_{ind}$	<b>Result</b>	$p_{cc}$	<b>Result</b>
HS	99%	0.040	<b>Fail</b>	0.081	Pass	0.026	<b>Fail</b>
	95%	0.858	Pass	0.030	<b>Fail</b>	0.093	Pass
EWMA	99%	0.533	Pass	0.466	Pass	0.631	Pass
	95%	0.858	Pass	0.671	Pass	0.899	Pass
GARCH	99%	0.063	Pass	0.357	Pass	0.116	Pass
	95%	0.980	Pass	0.106	Pass	0.271	Pass

As shown in Table 4, HS fails the unconditional and joint tests at 99% and the independence test at 95%, confirming it is inadequate for tail risk estimation. Both FHS methods pass all tests, but EWMA shows significantly higher p-values, indicating a robust fit.

**Table 5:**  
*ES backtesting results*

<b>Model</b>	<b>Level</b>	<b>p-value</b>	<b>Conclusion</b>
HS	99%	0.038	<b>Fail</b>
	95%	0.130	Pass
EWMA	99%	0.797	Pass
	95%	0.403	Pass
GARCH	99%	0.621	Pass
	95%	0.209	Pass

Table 5 reinforces that HS underestimates Expected Shortfall at extreme quantiles. Both FHS models pass, with EWMA again providing higher confidence.

**Recommendation:** We recommend the FHS-EWMA approach. It passes all backtests with the highest statistical confidence, accurately captures volatility clustering, and avoids the parameter instability observed in the rolling GARCH model. HS should be rejected due to its failure to adapt to market shocks and poor backtesting performance.

## References

- Bloomberg. (2025). About euro stoxx 50 price eur [Accessed: 2025-11-20]. <https://www.bloomberg.com/quote/SX5E:IND>
- Kupiec, P. H., et al. (1995). *Techniques for verifying the accuracy of risk measurement models* (Vol. 95). Division of Research; Statistics, Division of Monetary Affairs, Federal ...
- Nelson, L. S. (1998). The anderson-darling test for normality. *Journal of Quality Technology*, 30(3), 298–299.
- ProShares. (2025). Dog — short dow30 etf – proshares [Accessed: 2025-11-20]. <https://www.proshares.com/our-etfs/leveraged-and-inverse/dog>
- Scholz, F. W., & Stephens, M. A. (1987). K-sample anderson–darling tests. *Journal of the American Statistical Association*, 82(399), 918–924.

## A Additional Tables and Figures

**Table 6:**

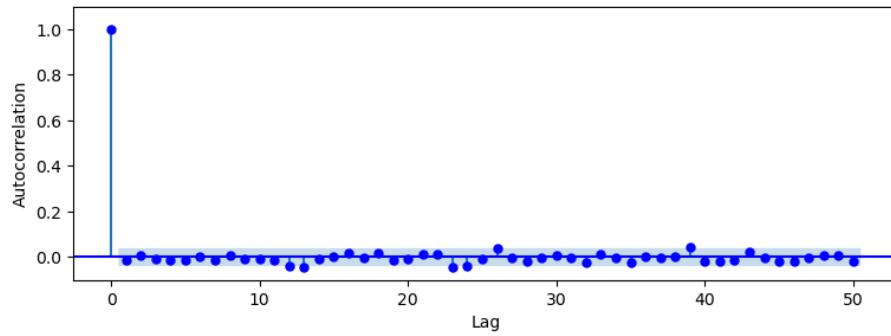
Coefficient estimates of the GARCH(1,1) model with constant mean and Student-t innovations

Variable	Estimate	Std. Error	T Value	$\mathbb{P} >  t $
$\mu$	0.0681	$1.72 \times 10^{-2}$	3.97	$7.22 \times 10^{-5}$
$\omega$	0.0467	$1.41 \times 10^{-2}$	3.31	$9.24 \times 10^{-4}$
$\alpha_1$	0.146	$2.78 \times 10^{-2}$	5.24	$1.61 \times 10^{-7}$
$\beta_1$	0.835	$2.89 \times 10^{-2}$	28.9	$3.12 \times 10^{-183}$
$\nu$	4.85	$4.65 \times 10^{-1}$	10.4	$2.00 \times 10^{-25}$

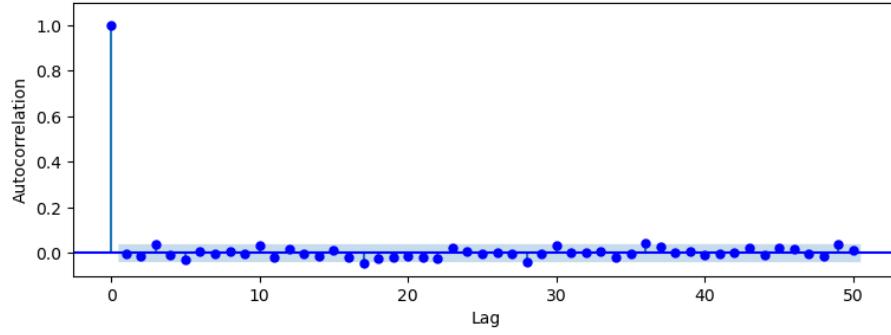
**Figure 13:**

Empirical ACFs of the standardised residuals and their absolute values the GARCH(1,1) model with constant mean and Student-t innovations

(a) Empirical ACF of the standardised residuals

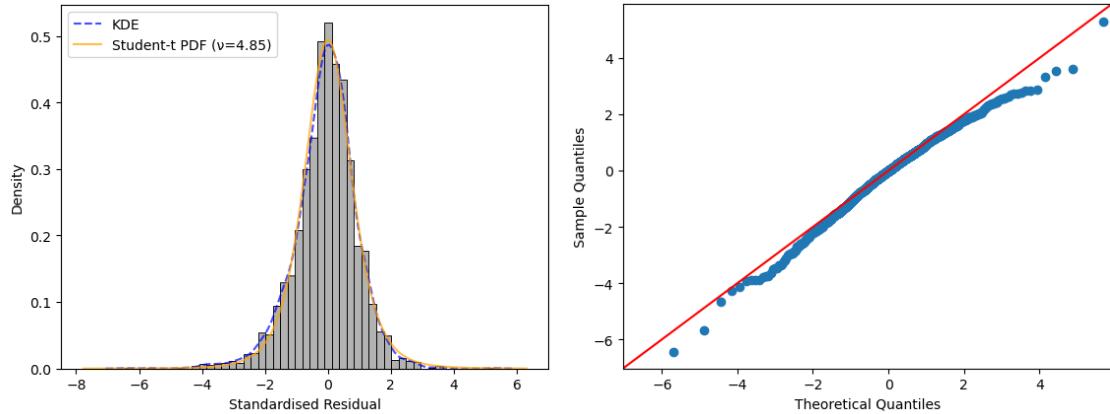


(b) Empirical ACF of the absolute standardised residuals



**Figure 14:**  
*Diagnostic plots of standardised residuals of the GARCH(1,1) model with constant mean and Student-t innovations*

(a) Histogram of the standardised residuals      (b) QQ-plot of standardised residuals



## B Code

```

1 import numpy as np
2 import pandas as pd
3 import matplotlib.pyplot as plt
4 import seaborn as sns
5 from scipy.stats import norm, t
6 from scipy import stats
7 import statsmodels.api as sm
8 from arch.univariate import arch_model
9
10 # Part A
11 # Compute log returns (%)
12 df = pd.read_csv('../data/QRM-2025-cw1-data-a.csv')
13 df_nodelay = df['Adj Close']
14 df_delay = df['Adj Close'].shift(1)
15 df['log_return'] = (np.log(df_nodelay) - np.log(df_delay)) * 100
16
17 # Part (i):
18 mu = df['log_return'].mean()
19 std = df['log_return'].std()
20 s = df['log_return'].skew()
21 k = df['log_return'].kurtosis()
22 print(f'log_return_mean: {mu:.5f}')
23 print(f'log_return_std: {std:.5f}')
24 print(f'log_return_skewness: {s:.5f}')
25 print(f'log_return_kurtosis: {k:.5f}')
26 print()
27 print(df['log_return'].describe())
28
29 # Plot histogram of log returns, KDE, and the fitted Gaussian pdf
30 x = np.linspace(-15,15,1000)
31 ax = sns.histplot(data=df['log_return'], color="gray", alpha = 0.6, bins=50, stat="density")
32 sns.kdeplot(data=df['log_return'], color="blue", alpha = 0.7, linestyle="--", label = 'KDE')
33 ax.plot(x, norm.pdf(x,mu,std),color = "orange", alpha = 0.7, label = 'Gaussian')
34 ax.set(xlabel='Log Return (%)', ylabel='Density')
35 ax.set_xlim(-15, 15)
36 ax.legend(loc = "upper left")

```

```

37 plt.show()
38
39 # Part (ii):
40 # Plot ACF of log returns, absolute log returns, and squared log returns
41 fig, ax = plt.subplots(figsize=(9, 3))
42 sm.graphics.tsa.plot_acf(
43     df['log_return'].dropna(),
44     lags=50,
45     ax=ax,
46     color="blue",
47 )
48 ax.set_title("")
49 ax.set_xlabel("Lag")
50 ax.set_ylabel("Autocorrelation")
51 ax.set_ylim(-0.1, 1.1)
52 plt.show()
53
54 fig, ax = plt.subplots(figsize=(9, 3))
55 sm.graphics.tsa.plot_acf(
56     np.abs(df['log_return']).dropna(),
57     lags=50,
58     ax=ax,
59     color="blue",
60 )
61 ax.set_title("")
62 ax.set_xlabel("Lag")
63 ax.set_ylabel("Autocorrelation")
64 ax.set_ylim(-0.1, 1.1)
65 plt.show()
66
67 fig, ax = plt.subplots(figsize=(9, 3))
68 sm.graphics.tsa.plot_acf(
69     (df['log_return'] ** 2).dropna(),
70     lags=50,
71     ax=ax,
72     color="blue",
73 )
74 ax.set_title("")
75 ax.set_xlabel("Lag")
76 ax.set_ylabel("Autocorrelation")
77 ax.set_ylim(-0.1, 1.1)
78 plt.show()
79
80 # Part (iii):
81 model = arch_model(df['log_return'].dropna(), mean='constant', p=1, q=1, dist='normal')
82 res = model.fit(update_freq = 3)
83 res.summary()
84
85 # Plot absolute log returns and fitted volatilities
86 df['Date'] = pd.to_datetime(df['Date'], dayfirst=True)
87 df2 = df.dropna(subset=['log_return']).copy()
88 abs_ret = df2['log_return'].abs()
89 vol = pd.Series(res.conditional_volatility, index=df2.index)
90 dates = df2['Date']
91
92 fig, ax = plt.subplots(figsize=(10,5))
93 ax.plot(dates, abs_ret, label='Absolute Log Returns (%)', alpha=0.7, color = 'orange')
94 ax.plot(dates, vol, label='Fitted Volatility (%)', linewidth=1.5, alpha=0.7, color = 'blue')
95 ax.set_xlabel('Date')
96 ax.set_ylabel('Value (%)')
97 ax.legend(loc = "upper left")
98 plt.show()
99
100 # Part (iv):
101 std_resid = res.std_resid
102

```

```

103 # Plot ACFs of standardised residuals and their absolute values
104 fig, ax = plt.subplots(figsize=(9, 3))
105 sm.graphics.tsa.plot_acf(
106     std_resid.dropna(),
107     lags=50,
108     ax=ax,
109     color="blue",
110 )
111 ax.set_title("")
112 ax.set_xlabel("Lag")
113 ax.set_ylabel("Autocorrelation")
114 ax.set_ylim(-0.1, 1.1)
115 plt.show()
116
117 fig, ax = plt.subplots(figsize=(9, 3))
118 sm.graphics.tsa.plot_acf(
119     np.abs(std_resid).dropna(),
120     lags=50,
121     ax=ax,
122     color="blue",
123 )
124 ax.set_title("")
125 ax.set_xlabel("Lag")
126 ax.set_ylabel("Autocorrelation")
127 ax.set_ylim(-0.1, 1.1)
128
129 # Plot histogram of standardised residuals
130 x = np.linspace(-4, 4, 200)
131 ax = sns.histplot(std_resid, color="gray", alpha = 0.6, bins=50, stat="density")
132 sns.kdeplot(data=std_resid, color="blue", alpha = 0.7, linestyle="--", label = 'KDE')
133 ax.plot(x,stats.norm.pdf(x),color = "orange", alpha = 0.7, label = 'Standard Normal PDF')
134 ax.set(xlabel='Standardised Residual', ylabel='Density')
135 plt.legend(loc = "upper left")
136 plt.show()
137
138 # Plot QQ-plot
139 sm.qqplot(std_resid, line='45', fit=True)
140 plt.show()
141
142 # A-D test
143 ad_res = stats.anderson(std_resid, dist = 'norm')
144 print(f'Anderson-Darling test:\n'
145       f'statistic = {ad_res.statistic:.1f}\n'
146       f'critical value with significance level 1% = {ad_res.critical_values[-1]:.2f}\n')
147
148 # Part (v)
149 df3 = df['log_return'].dropna()
150
151 def arma11_garch11_t_loglikelihood(alpha0, alphal, betal, arma_mu, phil, thetal, nu, x):
152
153     # Store time series of sigma_t^2 and mu_t
154     var = np.zeros_like(x)
155     mu = np.zeros_like(x)
156
157     # Set up proxies for sigma_0^2 and \mu_0
158     var[0] = np.var(x, ddof=1) # sample variance
159     mu[0] = np.mean(x) # sample mean
160
161     # Iterate over the ARMA(1,1)-GARCH(1,1) dynamics
162     for i in range(1, len(x)):
163         mu[i] = arma_mu + phil * (x.iloc[i-1] - arma_mu) + thetal * (x.iloc[i-1] - mu[i-1])
164         var[i] = alpha0 + alphal * (x.iloc[i-1] - mu[i-1])**2 + betal * var[i-1]
165
166     return np.sum(-0.5*np.log(var[1:]) + np.log(t.pdf((x.iloc[1:]-mu[1:])/np.sqrt(var[1:])), nu, 0, np.sqrt((nu-2)/nu))))

```

```

167
168 from scipy.optimize import minimize
169 from statsmodels.tsa.arima.model import ARIMA
170
171 objfun_armagarch_t = lambda y : -1 * armall_garch11_t_loglikelihood(y[0], y[1], y[2], y[3], y[4], y[5], y[6], df3)
172 constraint = lambda y : 1 - y[1] - y[2]
173
174 uncond_var = np.var(df3, ddof=1)
175 alpha1_0 = 0.1
176 betal_0 = 0.85
177 alpha0_0 = uncond_var * (1 - alpha1_0 - betal_0)
178
179 arma_0 = ARIMA(df3, order=(1, 0, 1)).fit()
180
181 arma_mu_0 = df3.mean()
182 phil_0 = arma_0.params.get("ar.L1", 0.0)
183 thetal_0 = arma_0.params.get("ma.L1", 0.0)
184
185 y0 = [alpha0_0, alpha1_0, betal_0, arma_mu_0, phil_0, thetal_0, 10]
186 print("Initial objective function value: " + str(objfun_armagarch_t(y0)))
187
188 bnds = ((0,np.inf), (0,1), (0,1), (-np.inf, np.inf), (-np.inf, np.inf), (-np.inf, np.inf),
189 ), (2.01, np.inf))
190
191 con = {'type': 'ineq', 'fun': constraint}
192
193 sol = minimize(objfun_armagarch_t,y0,method='SLSQP',bounds=bnds,constraints=con, options
194 ={'disp': True})
195
196 # Coefficient estimates
197 print("Estimate of alpha0: " + str(sol.x[0]))
198 print("Estimate of alpha1: " + str(sol.x[1]))
199 print("Estimate of betal: " + str(sol.x[2]))
200 print("Estimate of mu: " + str(sol.x[3]))
201 print("Estimate of phil: " + str(sol.x[4]))
202 print("Estimate of thetal: " + str(sol.x[5]))
203 print("Estimate of nu: " + str(sol.x[6]))
204
205 # AIC and BIC
206 k = len(sol.x)
207 n = len(df3)
208 ll = -objfun_armagarch_t(sol.x)
209 AIC = 2*k - 2*ll
210 BIC = k*np.log(n) - 2*ll
211 print(AIC, BIC)
212
213 # Standard Errors
214 from numdifftools import Hessian
215
216 # Evaluate Hessian of NEGATIVE loglikelihood (objective function)
217 H = Hessian(objfun_armagarch_t)(sol.x)
218 # Inverse Hessian = covariance matrix
219 cov = np.linalg.inv(H)
220 # Standard errors
221 std_err = np.sqrt(np.diag(cov))
222 print(std_err)
223
224 # T-values
225 t_values = sol.x / std_err
226 print(t_values)
227
228 # 95% CIs
229 ci_lower = sol.x - 1.96 * std_err
230 ci_upper = sol.x + 1.96 * std_err
231 print(ci_lower,
232 ci_upper)

```

```

231
232 # Calculate fitted values and residuals
233 alpha0_hat = sol.x[0]
234 alpha1_hat = sol.x[1]
235 betal_hat = sol.x[2]
236 arma_mu_hat = sol.x[3]
237 phil_hat = sol.x[4]
238 thetal_hat = sol.x[5]
239 nu_hat = sol.x[6]
240
241 def armall_garch11_paths(params, x):
242     alpha0, alphal, betal, arma_mu, phil, thetal, nu = params
243
244     x = x.to_numpy() # work with ndarray for speed
245     n = len(x)
246
247     var = np.empty(n)
248     mu = np.empty(n)
249
250     # Initial values
251     var[0] = np.var(x, ddof=1)
252     mu[0] = np.mean(x)
253
254     for i in range(1, n):
255         # Mean recursion
256         mu[i] = arma_mu + phil * (x[i-1] - arma_mu) + thetal * (x[i-1] - mu[i-1])
257         # Variance recursion
258         e_prev = x[i-1] - mu[i-1]
259         var[i] = alpha0 + alphal * e_prev * e_prev + betal * var[i-1]
260
261     # Standardized residuals
262     sigma = np.sqrt(var)
263     z = (x - mu) / sigma
264
265     # Wrap back into Series with same index as x
266     idx = df3.index # or x_index saved separately if you prefer
267     mu_s = pd.Series(mu, index=idx, name="mu_hat")
268     var_s = pd.Series(var, index=idx, name="var_hat")
269     sigma_s = pd.Series(sigma, index=idx, name="sigma_hat")
270     z_s = pd.Series(z, index=idx, name="z_hat")
271
272     return mu_s, var_s, sigma_s, z_s
273
274 params_hat = (alpha0_hat, alphal_hat, betal_hat,
275                 arma_mu_hat, phil_hat, thetal_hat, nu_hat)
276 mu_hat, var_hat, sigma_hat, z_hat = armall_garch11_paths(params_hat, df3)
277 z_std = z_hat.iloc[1:]
278 z = z_std.dropna()
279
280 # X-grid for plotting the fitted t-density
281 x = np.linspace(z.min() * 1.2, z.max() * 1.2, 200)
282 # Scale parameter to ensure Var = 1
283 scale_hat = np.sqrt((nu_hat - 2) / nu_hat)
284
285 # Histogram of standardised residuals
286 ax = sns.histplot(z, color="gray", alpha=0.6, bins=50, stat="density")
287 # KDE
288 sns.kdeplot(
289     data=z,
290     color="blue",
291     alpha=0.7,
292     linestyle="--",
293     label="KDE"
294 )
295 # Fitted Student-t PDF
296 ax.plot(
297     x,

```

```

298     t.pdf(x, df=nu_hat, loc=0, scale=scale_hat),
299     color="orange",
300     alpha=0.7,
301     label=f"Student-t PDF"
302 )
303 ax.set(xlabel="Standardised Residual", ylabel="Density")
304 plt.legend(loc="upper left")
305 plt.show()
306
307 # QQ-plot
308 sm.qqplot(
309     z,
310     dist=t,
311     distargs=(nu_hat,),
312     loc=0,
313     scale=scale_hat,
314     line='45'
315 )
316 plt.show()
317
318 # ACFs of standardised residuals
319 fig, ax = plt.subplots(figsize=(9, 3))
320 sm.graphics.tsa.plot_acf(
321     z.dropna(),
322     lags=50,
323     ax=ax,
324     color="blue",
325 )
326 ax.set_title("")
327 ax.set_xlabel("Lag")
328 ax.set_ylabel("Autocorrelation")
329 ax.set_ylim(-0.1, 1.1)
330
331 fig, ax = plt.subplots(figsize=(9, 3))
332 sm.graphics.tsa.plot_acf(
333     np.abs(z).dropna(),
334     lags=50,
335     ax=ax,
336     color="blue",
337 )
338 ax.set_title("")
339 ax.set_xlabel("Lag")
340 ax.set_ylabel("Autocorrelation")
341 ax.set_ylim(-0.1, 1.1)
342
343 from scipy.stats import anderson_ksamp
344 # Simulate reference t-distribution
345 t_sim = t.rvs(df=nu_hat, loc=0, scale=scale_hat, size=50000, random_state=42)
346 # A-D k-sample test
347 ad_stat, crit_vals, sig_level = anderson_ksamp([z, t_sim])
348
349 # Without ARMA(1,1)
350 model2 = arch_model(df['log_return'].dropna(), mean='constant', p=1, q=1, dist='t')
351 res2 = model2.fit(update_freq = 3)
352 res2.summary()
353
354 # Standardised residuals and diagnostic plots
355 std_resid2 = res2.std_resid
356 nu = res2.params["nu"]
357
358 # X-grid for the theoretical t density
359 x_vals = np.linspace(std_resid2.min() * 1.2, std_resid2.max() * 1.2, 500)
360 # Scale so that Var = 1 (same as in your likelihood)
361 scale2 = np.sqrt((nu - 2.0) / nu)
362
363 # Histogram
364 ax = sns.histplot(std_resid2, color="gray", alpha=0.6, bins=50, stat="density")

```

```

365 # KDE
366 sns.kdeplot(
367     data=std_resid2,
368     color="blue",
369     alpha=0.7,
370     linestyle="--",
371     label="KDE"
372 )
373 # Fitted Student-t PDF
374 ax.plot(
375     x,
376     t.pdf(x, df=nu, loc=0, scale=scale2),
377     color="orange",
378     alpha=0.7,
379     label=f"Student-t PDF"
380 )
381 ax.set(xlabel="Standardised Residual", ylabel="Density")
382 plt.legend(loc="upper left")
383 plt.show()
384
385 # QQ-plots
386 sm.qqplot(
387     std_resid2,
388     dist=t,
389     distargs=(nu,),    # only df here
390     loc=0,
391     scale=scale2,
392     line='45'
393 )
394 plt.show()
395
396 # ACFs of standardised residuals
397 fig, ax = plt.subplots(figsize=(9, 3))
398 sm.graphics.tsa.plot_acf(
399     std_resid2.dropna(),
400     lags=50,
401     ax=ax,
402     color="blue",
403 )
404 ax.set_title("")
405 ax.set_xlabel("Lag")
406 ax.set_ylabel("Autocorrelation")
407 ax.set_ylim(-0.1, 1.1)
408
409 fig, ax = plt.subplots(figsize=(9, 3))
410 sm.graphics.tsa.plot_acf(
411     np.abs(std_resid2).dropna(),
412     lags=50,
413     ax=ax,
414     color="blue",
415 )
416 ax.set_title("")
417 ax.set_xlabel("Lag")
418 ax.set_ylabel("Autocorrelation")
419 ax.set_ylim(-0.1, 1.1)
420
421 # Simulate reference t-distribution
422 t_sim2 = t.rvs(df=nu, loc=0, scale=scale2, size=50000, random_state=42)
423 # A-D ksamp test
424 ad_stat2, crit_vals2, sig_level2 = anderson_ksamp([std_resid2, t_sim2])
425
426 # Part B
427 import pandas as pd
428 import numpy as np
429 import matplotlib.pyplot as plt
430 from numpy.lib.stride_tricks import sliding_window_view
431 from arch import arch_model

```

```

432 from scipy.stats import chi2
433 from scipy.stats import norm
434 import seaborn as sns
435
436 # Initialisation
437 df = pd.read_csv('../data/QRM-2025-cwl-data-b.csv')
438 prices = df['Adj Close'].values
439 dates = pd.to_datetime(df['Date'], dayfirst=True).values[1:]
440 losses = -np.log(prices[1:] / prices[:-1])
441 T = len(losses)
442
443 # Common HS Setup
444 N = 500 # Window Size
445 targets = losses[N:]
446 target_dates = dates[N:]
447
448 # Part (i):
449 windows = sliding_window_view(losses[:-1], window_shape=N) # Technical reason to exclude
450         the last date
451 var_95_i = np.quantile(windows, 0.95, axis=1)
452 es_95_i = np.array([w[w >= v].mean() for w, v in zip(windows, var_95_i)])
453 var_99_i = np.quantile(windows, 0.99, axis=1)
454 es_99_i = np.array([w[w >= v].mean() for w, v in zip(windows, var_99_i)])
455
456 plt.figure(figsize=(12, 6))
457 plt.plot(target_dates, targets, 'gray', alpha=0.6, lw=0.8, label='Loss')
458 plt.plot(target_dates, var_95_i, 'orange', linestyle='--', label='95% VaR')
459 plt.plot(target_dates, es_95_i, 'orange', alpha=0.7, label='95% ES')
460 plt.plot(target_dates, var_99_i, 'blue', linestyle='--', label='99% VaR')
461 plt.plot(target_dates, es_99_i, 'blue', alpha=0.7, label='99% ES')
462 plt.ylim(bottom=0) # Only include positive losses
463 plt.title(f'HS Forecasts (N={N})')
464 plt.ylabel('Linearised Loss Relative to the Value')
465 plt.legend(loc='upper left')
466 plt.grid(True, alpha=0.3)
467 plt.savefig('assets/hs.png')
468 plt.show()
469
470 # Part (ii)
471 sigma2 = np.zeros(T)
472 sigma2[0] = np.var(losses[:N]) # Take variance as initial condition (DUE TO DISCUSS)
473
474 # Run EWMA
475 for t in range(1, T):
476     sigma2[t] = (1 - 0.06) * sigma2[t-1] + 0.06 * losses[t-1]**2
477 sigma = np.sqrt(sigma2)
478 plt.figure(figsize=(12, 6))
479 plt.plot(dates, sigma2, 'orange', alpha=0.6, lw=0.8, label='Volatility based on EWMA
480         method')
481 plt.title('EWMA Volatility Forecasts')
482 plt.legend(loc='upper left')
483 plt.grid(True, alpha=0.3)
484 plt.savefig('assets/ewma_vol.png')
485 plt.show()
486
487 std_residuals = losses / sigma # Assume mean is zero (DUE TO DISCUSS)
488 z_window = std_residuals[-500:]
489
490 plt.figure(figsize=(8, 4))
491 ax = sns.histplot(data = z_window, bins=50, stat = "density", color = 'gray', alpha =
492         0.6, label='Empirical $F_Z$ (last 500 days)')
493 xs = np.linspace(z_window.min() - 1, z_window.max() + 1, 200)
494 plt.plot(xs, norm.pdf(xs), linewidth=2, linestyle = '--',
495         label='$N(0,1)$')
496 plt.title('Histogram of Standardised EWMA Residuals (Approx. $F_Z$)')
497 plt.xlabel('Standardised residual $z$')

```

```

496 plt.ylabel('Density')
497 plt.legend()
498 plt.tight_layout()
499 plt.savefig('assets/Distribution_Fz_EWMA.png')
500 plt.show()
501
502 windows = sliding_window_view(std_residuals[:-1], window_shape=N)
503 vols = sigma[N:]
504
505 var_95_z_ii = np.quantile(windows, 0.95, axis=1)
506 es_95_z_ii = np.array([w[w >= v].mean() for w, v in zip(windows, var_95_z_ii)])
507 var_95_ii = vols * var_95_z_ii
508 es_95_ii = vols * es_95_z_ii
509
510 var_99_z_ii = np.quantile(windows, 0.99, axis=1)
511 es_99_z_ii = np.array([w[w >= v].mean() for w, v in zip(windows, var_99_z_ii)])
512 var_99_ii = vols * var_99_z_ii
513 es_99_ii = vols * es_99_z_ii
514
515 plt.figure(figsize=(12, 6))
516 plt.plot(target_dates, targets, 'gray', alpha=0.6, lw=0.8, label='Loss')
517 plt.plot(target_dates, var_95_ii, 'orange', linestyle='--', label='95% VaR')
518 plt.plot(target_dates, es_95_ii, 'orange', alpha=0.7, label='95% ES')
519 plt.plot(target_dates, var_99_ii, 'blue', linestyle='--', label='99% VaR')
520 plt.plot(target_dates, es_99_ii, 'blue', alpha=0.7, label='99% ES')
521 plt.ylim(bottom=0) # Only include positive losses
522 plt.title(f'FHS-EWMA Forecasts (N={N})')
523 plt.ylabel('Linearised Loss Relative to the Value')
524 plt.legend(loc='upper left')
525 plt.grid(True, alpha=0.3)
526 plt.savefig('assets/fhs_ewma.png')
527 plt.show()
528
529 # Part (iii)
530 windows = sliding_window_view(losses[:-1], window_shape=N)
531 var_95_iii = np.zeros(T-N)
532 es_95_iii = np.zeros(T-N)
533 var_99_iii = np.zeros(T-N)
534 es_99_iii = np.zeros(T-N)
535 for i in range(T-N):
536     am = arch_model(windows[i] * 100, mean='Constant', vol='Garch', p=1, q=1, dist='normal')
537     res = am.fit(update_freq=0, disp='off')
538
539     std_residuals = res.std_resid
540
541     var_95_z_iii = np.quantile(std_residuals, 0.95)
542     es_95_z_iii = std_residuals[std_residuals >= var_95_z_iii].mean()
543
544     var_99_z_iii = np.quantile(std_residuals, 0.99)
545     es_99_z_iii = std_residuals[std_residuals >= var_99_z_iii].mean()
546
547     forecast = res.forecast(horizon=1)
548
549     mu_forecast = forecast.mean.iloc[-1].values[0] / 100.0
550     sigma_forecast = np.sqrt(forecast.variance.iloc[-1].values[0]) / 100.0
551
552     var_95_iii[i] = mu_forecast + sigma_forecast * var_95_z_iii
553     es_95_iii[i] = mu_forecast + sigma_forecast * es_95_z_iii
554
555     var_99_iii[i] = mu_forecast + sigma_forecast * var_99_z_iii
556     es_99_iii[i] = mu_forecast + sigma_forecast * es_99_z_iii
557
558 last_window = losses[-N:] * 100
559 am_last = arch_model(last_window, mean='Constant', vol='Garch',
560                      p=1, q=1, dist='normal')
561 res_last = am_last.fit(update_freq=0, disp='off')

```

```

562 z_last = res_last.std_resid
563
564 plt.figure(figsize=(8, 4))
565 sns.histplot(z_last, bins=50, stat="density",
566               color="gray")
567 sns.kdeplot(z_last, color="#1f77b4", linestyle="--", linewidth=2,
568               label="KDE")
569
570 x = np.linspace(z_last.min()*1.2, z_last.max()*1.2, 400)
571 plt.plot(x, norm.pdf(x), color="#ff7f0e", linewidth=2, label="Gaussian")
572
573 plt.title("Histogram of Standardised GARCH Residuals (Approx. $F_Z$)")
574 plt.xlabel("Standardised residual $z$")
575 plt.ylabel("Density")
576 plt.legend()
577 plt.tight_layout()
578 plt.savefig('assets/Distribution_Fz_GARCH.png')
579 plt.show()
580
581 plt.figure(figsize=(12, 6))
582 plt.plot(target_dates, targets, 'gray', alpha=0.6, lw=0.8, label='Loss')
583 plt.plot(target_dates, var_95_iii, 'orange', linestyle='--', label='95% VaR')
584 plt.plot(target_dates, es_95_iii, 'orange', alpha=0.7, label='95% ES')
585 plt.plot(target_dates, var_99_iii, 'blue', linestyle='--', label='99% VaR')
586 plt.plot(target_dates, es_99_iii, 'blue', alpha=0.7, label='99% ES')
587 plt.ylim(bottom=0) # Only include positive losses
588 plt.title(f'FHS-GARCH Forecasts (N={N})')
589 plt.ylabel('Linearised Loss Relative to the Value')
590 plt.legend(loc='upper left')
591 plt.grid(True, alpha=0.3)
592 plt.savefig('assets/fhs_garch.png')
593 plt.show()
594
595 # Backtesting
596 # Violations
597 v_95_i = targets > var_95_i
598 v_99_i = targets > var_99_i
599 v_95_ii = targets > var_95_ii
600 v_99_ii = targets > var_99_ii
601 v_95_iii = targets > var_95_iii
602 v_99_iii = targets > var_99_iii
603
604 def backtest(violations, alpha):
605
606     T = len(violations)
607     N1 = np.sum(violations)
608     N0 = T - N1
609     eps = 1e-15 # Prevent divide by zero
610
611     # Unconditional Coverage
612     pi = N1 / T
613
614     logL_H0 = N0 * np.log(alpha) + N1 * np.log(1-alpha)
615     logL_H1 = N0 * np.log(1 - pi) + N1 * np.log(pi)
616     LR_uc = -2 * (logL_H0 - logL_H1)
617     p_uc = 1 - chi2.cdf(LR_uc, 1)
618
619     # Independence
620     t00, t01, t10, t11 = 0, 0, 0, 0
621     for i in range(1, T):
622         if not violations[i-1] and not violations[i]: t00 += 1
623         if not violations[i-1] and violations[i]: t01 += 1
624         if violations[i-1] and not violations[i]: t10 += 1
625         if violations[i-1] and violations[i]: t11 += 1
626
627     pi01 = t01 / (t00 + t01)
628     pill = t11 / (t10 + t11)

```

```

629
630     logL_H0 = (t00 + t10) * np.log(1 - pi + eps) + (t01 + t11) * np.log(pi + eps)
631     logL_H1 = t00 * np.log(1 - pi01 + eps) + t01 * np.log(pi01 + eps) + t10 * np.log(1 -
632         pi11 + eps) + t11 * np.log(pi11 + eps)
633
634     LR_ind = -2 * (logL_H0 - logL_H1)
635     p_ind = 1 - chi2.cdf(LR_ind, 1)
636
637     # Joint
638     LR_cc = LR_uc + LR_ind
639     p_cc = 1.0 - chi2.cdf(LR_cc, 2)
640
641     return LR_uc, p_uc, LR_ind, p_ind, LR_cc, p_cc
642
643 models = [
644     ('HS', '95%', v_95_i, 0.95),
645     ('HS', '99%', v_99_i, 0.99),
646     ('FHS-EWMA', '95%', v_95_ii, 0.95),
647     ('FHS-EWMA', '99%', v_99_ii, 0.99),
648     ('FHS-GARCH', '95%', v_95_iii, 0.95),
649     ('FHS-GARCH', '99%', v_99_iii, 0.99),
650 ]
651 results_uc = []
652 results_ind = []
653 results_cc = []
654 for name, alpha_str, violations, alpha_val in models:
655     LR_uc, p_uc, LR_ind, p_ind, LR_cc, p_cc = backtest(violations, alpha_val)
656     results_uc.append({
657         'Model': name,
658         'alpha': alpha_str,
659         'LR_uc': LR_uc,
660         'p_uc': p_uc
661     })
662     results_ind.append({
663         'Model': name,
664         'alpha': alpha_str,
665         'LR_ind': LR_ind,
666         'p_ind': p_ind
667     })
668     results_cc.append({
669         'Model': name,
670         'alpha': alpha_str,
671         'LR_cc': LR_cc,
672         'p_cc': p_cc
673     })
674 df_results_uc = pd.DataFrame(results_uc)
675 df_results_ind = pd.DataFrame(results_ind)
676 df_results_cc = pd.DataFrame(results_cc)
677 print(df_results_uc.to_string(index=False))
678 print(df_results_ind.to_string(index=False))
679 print(df_results_cc.to_string(index=False))

680 # Function for backtest ES for different model
681 def es_backtest(loss, var, es):
682     #Indicator function for whether the loss exceeding the VaR
683     I = (loss > var).astype(float)
684     # value of loss compared to ES forecast
685     ex1 = (loss - es) * I
686     sum_ex1 = np.sum(ex1)
687     sum_ex2 = np.sum(ex1**2)
688     #Standard normal test statistics
689     Z = sum_ex1 / np.sqrt(sum_ex2)
690     # p-value
691     p = 1 - norm.cdf(Z)
692     return {'Test Statistics': Z, 'p-value': p}
693
694 # Create a data frame for backtest ES result

```

```
695 ES_backtest = pd.DataFrame({ "HS-ES-95": es_backtest(losses[500:],var_95_i,es_95_i),  
696             "HS-ES-99": es_backtest(losses[500:],var_99_i,es_99_i),  
697             "FHS-EWMA_ES-95": es_backtest(losses[500:],var_95_ii,es_95_ii),  
698             "FHS-EWMA_ES-99": es_backtest(losses[500:],var_99_ii,es_99_ii),  
699             "FHS-GARCH_ES-95": es_backtest(losses[500:],var_95_iii,es_95_iii),  
700             "FHS-GARCH_ES-99": es_backtest(losses[500:],var_99_iii,es_99_iii)  
701         }) .T  
701 print(ES_backtest)
```